

**Electronics and Computer Science**  
Faculty of Engineering and Physical Sciences  
University of Southampton

**Ashwinkrishna Azhagesh**

**25/03/2025**

**An AI Approach to Chaotic Physical Systems:**

Project supervisor: **Adam Peugeot**

Second examiner: **David Millard**

Progress report submitted for the award of  
**Bachelors of Science**

# Abstract

Empirical laws are mathematical generalisations found through observing the physical world. It has taken us centuries of gathering data, keen research along with repeated experiments, and no doubt plenty of talented scientists to discover these laws. Leading us to understand everything from the mysteries that govern the collision of two objects to the shape of the path planets thread upon.

Recent advances in neural networks including increases in computational power permit us to train models, that replicate, fasten and automate our discovery of empirical laws. This extends to even noisy chaotic systems such as the double pendulum. Combined with white box models, symbolic regression and explainable A.I., we can peer into the "mind," of how such models, process data and conclude their observations. Human cognition is inherently finite in its capacity for thought and observational ability, has been historically overcome through the development of new tools such as the microscope. Similarly, cognitive biases can be mitigated, by utilising artificial intelligence, which is a rapidly emerging technology capable of expanding our perception and analysis.

### **Statement of Originality**

- I have read and understood the ECS Academic Integrity information and the University's Academic Integrity Guidance for Students.

- I am aware that failure to act in accordance with the Regulations Governing Academic Integrity may lead to the imposition of penalties which, for the most serious cases, may include termination of programme.

- I consent to the University copying and distributing any or all of my work in any form and using third parties (who may be based outside the EU/EEA) to verify whether my work contains plagiarised material, and for quality assurance purposes.

**You must change the statements in the boxes if you do not agree with them.**

We expect you to acknowledge all sources of information (e.g. ideas, algorithms, data) using citations. You must also put quotation marks around any sections of text that you have copied without paraphrasing. If any figures or tables have been taken or modified from another source, you must explain this in the caption and cite the original source.

**I have acknowledged all sources, and identified any content taken from elsewhere.**

If you have used any code (e.g. open-source code), reference designs, or similar resources that have been produced by anyone else, you must list them in the box below. In the report, you must explain what was used and how it relates to the work you have done.

**I have not used any resources produced by anyone else.**

You can consult with module teaching staff/demonstrators, but you should not show anyone else your work (this includes uploading your work to publicly-accessible repositories e.g. Github, unless expressly permitted by the module leader), or help them to do theirs. For individual assignments, we expect you to work on your own. For group assignments, we expect that you work only with your allocated group. You must get permission in writing from the module teaching staff before you seek outside assistance, e.g. a proofreading service, and declare it here.

**I did all the work myself, or with my allocated group, and have not helped anyone else.**

We expect that you have not fabricated, modified or distorted any data, evidence, references, experimental results, or other material used or presented in the report. You must clearly describe your experiments and how the results were obtained, and include all data, source code and/or designs (either in the report, or submitted as a separate file) so that your results could be reproduced.

**The material in the report is genuine, and I have included all my data/code/designs.**

We expect that you have not previously submitted any part of this work for another assessment. You must get permission in writing from the module teaching staff before re-using any of your previously submitted work for

this assessment.

**I have not submitted any part of this work for another assessment.**

If your work involved research/studies (including surveys) on human participants, their cells or data, or on animals, you must have been granted ethical approval before the work was carried out, and any experiments must have followed these requirements. You must give details of this in the report, and list the ethical approval reference number(s) in the box below.

**My work did not involve human participants, their cells or data, or animals.**

ECS Statement of Originality Template, updated August 2018, Alex Weddell aiofficer@ecs.soton.ac.uk

# Acknowledgements

I would like to thank my supervisors, Professor Adam Peugeot and Professor David Millard, for all the help and advise I received throughout this project.

# Contents

<b>Abstract</b>	<b>8</b>
<b>Statement of Originality</b>	<b>8</b>
<b>1 Introduction:</b>	<b>8</b>
1.1 Motivation:	8
<b>2 Background Research:</b>	<b>9</b>
2.1 Literature Review:	9
2.2 Symbolic Regression:	9
2.2.1 Brute Force:	9
2.2.2 Partial Derivatives:	9
2.2.3 Dimensional Analysis:	11
2.2.4 Genetic Programming:	11
2.2.5 Broader Use Cases and Techniques:	12
<b>3 PySR</b>	<b>12</b>
3.1 Momentum Laws:	12
3.2 Pendulum Laws:	13
3.3 Noise:	14
3.3.1 How noise affects the model:	14
3.3.2 Denoise function:	16
<b>4 Symbolic Regression from First Principles:</b>	<b>17</b>
4.1 A Brute-Force Approach:	17
4.1.1 Exploiting Physical Properties:	18
4.1.2 Dealing with constants:	19
4.1.3 Dealing with powers:	19
4.1.4 Chaining powers and constants:	20
4.1.5 Loading data:	20
4.1.6 How to mitigate noise in data:	21
4.1.7 Evaluating expressions:	21
4.2 Polynomial Fit Module:	22
4.2.1 Data Loading:	23
4.2.2 Generating polynomial expressions:	23
4.2.3 Filtering the Polynomial expressions:	23
4.2.4 Evaluating expressions:	24
4.2.5 Best Polynomial Fit:	25
4.3 Dimensional Analysis:	25
4.3.1 Handling Units:	25
4.3.2 Construct Matrix and Target Vector:	26
4.3.3 Solving Dimension and Basis Units:	26
4.3.4 Data Transformation Function:	27
4.3.5 Symbolic Transformation Generator:	28
4.4 Neural Network Fitting:	29
4.4.1 Symbolic Network:	29
4.4.2 Preparing the data:	30
4.4.3 Training the Network:	31
4.4.4 Predict Function:	32
4.5 Pareto Frontier Optimisation:	33
4.5.1 Points:	33

4.5.2	Pareto Point Set: . . . . .	34
4.6	Plotting: . . . . .	35
4.6.1	Plot for Pareto front: . . . . .	35
4.7	Main method bringing it together Regressor: . . . . .	36
4.7.1	main solver . . . . .	36
4.8	Testing: . . . . .	37
4.8.1	Unit Testing: . . . . .	37
4.8.2	Integration Testing: . . . . .	38
4.8.3	Performance Testing: . . . . .	38
<b>5</b>	<b>Applying the model to Biological Data:</b>	<b>39</b>
<b>6</b>	<b>Evaluation:</b>	<b>39</b>
	<b>References</b>	<b>40</b>
	<b>Appendix:</b>	<b>41</b>
<b>7</b>	<b>Project Management:</b>	<b>41</b>
7.1	Risk Assessment: . . . . .	41
7.2	Gantt Charts: . . . . .	42
7.3	Project Reflection: . . . . .	42
<b>8</b>	<b>Conclusion:</b>	<b>44</b>

# 1 Introduction:

## 1.1 Motivation:

Humanity has spent millennia observing the world, creating concepts that describe the variables, such as mass and force, to derive laws. In physics, like with all human endeavours, new discoveries and ways of thought are based upon previous works, creating a natural bias in the way new problems are approached. All existing theories, are therefore somewhat biased, this combined with our pre-existing bias in our biological brains, can introduce some hurdles to our future progress [26] [18].

In the 17th Century, Kepler had gotten his hands on the world's most precise data tables on the orbits on planets, using this he spent close to half a decade, and after numerous unsuccessful attempts, he had began a scientific revolution at the time, describing Mar's orbit to be an ellipse [13]. In essence, scientists throughout history, much like Kepler, have spent a great deal of time, discovering the right expressions to match the relevant data they have, this at it's core is symbolic regression. Now, a few centuries later, even with exponential increases in orders of magnitude in our capability to perform calculations through computers, the process of discovering natural laws and the way to express them, has to some extent resisted automation.

One of the core challenges of physics and artificial intelligence, is finding analytical relations automatically, discovering a symbolic expression that accurately matches the data from an unknown function. This problem, due to it's nature, is NP-hard [11] in principle. The vastness of the space of mathematical constants, adds to the difficulty. This aligns with a growing research area focused on integrating scientific domain knowledge with machine learning techniques to accelerate discovery and overcome limitations of purely data-driven or purely mechanistic models [25]. This literature review aims to present the recent advances in discovery of empirical laws through data powered by artificial intelligence. It focuses on methodologies that diminish human bias through seeking solutions without assumptions. We will explore various techniques employed to achieve these goals, which includes reducing the search space, and analyse the effectiveness of these methods.

As there is more data being generated than ever before and new experiments, we need a systematic and automatic way to deduce various mathematical patterns and laws in these data. Through the use of symbolic regression we can utilise these data, and in an explainable manner deduce various new physical laws. In this research I have also extended this beyond physics and have applied this to biological data sets which is a novel application of this method.




Figure 1: This is the orbit of Earth and Mars around the Sun.



## 2 Background Research:

### 2.1 Literature Review:

### 2.2 Symbolic Regression:

Symbolic regression, is a technique that analyses and searches over the space of traceable mathematical expressions to find the best fit for a data set. By not requiring prior information about the model, it is unbiased. There are a plethora of various strategies [8] that have been implemented in solving for empirical laws [19], we will explore some of them below. It is also worth mentioning, that unlike other well-known techniques for regression, (eg: neural networks), that are essentially black boxes [2], symbolic regression, aims to extract white-box models and is easy to analyse.

#### 2.2.1 Brute Force:

Symbolic Regression (SR), is interpretable [3], unlike Neural Networks (NN), which are often considered more explainable. The difference is interpretability allows us to comprehend how the model works, like observing how gears move in a glass box, while explainable means you get an overview of why a certain output was achieved, even without knowing the full nuances of it's inner workings. This however does not mean there aren't ways to approach this incrementally with neural networks, which allows for interpretability to increase [?].

There however, are some challenges associated with SR, in comparison to function fitting (NN). SR, starts with nothing, a blank slate, and it has to learn the entire expression [10], unlike function fitting which just tweaks an already existing function. The exponential search space [27], causes it to be extremely computationally expensive to explore all possibilities. This combined with the fact that, most optimisation algorithms expect a smooth search space [16], however SR lacks smooth interpolation, small changes in the potential solutions (expression), ie:  $x^3 + 0.1$  can significantly alter the output. Finally, if the nature of the problem is badly posed [17], there might potentially be multiple solutions to the same data. Imagine trying to find a single polynomial equation with only two points of data, the need to balance finding accurate expressions with finding the most simplistic and generalisable fit, is sometimes troublesome.

The brute force approach of simply trying all possible combinations of symbolic expressions within some defined space. The model will subsequently increase the complexity over time, and will stop when either the fitting errors lowers below some defined limit or exceeds the upper limit of runtime. While in theory can solve all of our problems, in practise takes longer than the age of our universe to finish. In essence it's like searching for a singular drop in the ocean. Thankfully, there are some ways of pruning the search space, and drastically reducing the time taken to solve for the most accurate expression.

#### 2.2.2 Partial Derivatives:

Partial derivatives, of some function  $f$ , with multiple variables such as  $x$  and  $y$ , is it's derivative with respect to one of those two variables, while the other variables in the function are kept constant. Formally, given a function with two or more variables,  $f(x_1, x_2, \dots, x_n)$ , the partial derivative of  $f$  with respect to  $x_i$ , where  $x_i$  is some value  $x$  in  $(x_1, x_2, \dots, x_i, \dots, x_n)$ , gives the rate of change of  $f$  with respect to  $x_i$ . It is calculated by taking the  $i$ th derivative of  $f$  with respect to  $x_i$ , whilst holding the other variables fixed. [21] [20]

The partial derivative of a function  $f(x, y)$  with respect to  $x$  is denoted  $\frac{\partial f}{\partial x}$  [12] and is defined:

$$\frac{\partial f}{\partial x} = \lim_{h \rightarrow 0} \left[ \frac{f(x+h, y) - f(x, y)}{h} \right]$$

Once you pass in the experimental data, you can pre-process the data, using calculated partial derivatives, for every pair of existing variables. Many physical laws, involve rates of change, and partial derivatives help us represent them. Furthermore it also guides the search process, as the algorithm can use the derivative to accurately represent the underlying laws involved. Through comparing how well the partial derivatives derived through the experimental data compared to the potential expression, the algorithm can assess the accuracy and feasibility of the expressions involved. This strategy can even be extended to prune the search space further, this could be achieved through incorporating knowledge of physics into the constraints for the partial derivatives. These concepts will be illustrated with an example below.




Figure 2: Temperature distribution along a heated rod, illustrating the physical scenario where partial derivatives describe heat flow.

Consider a iron rod, that has been heated up, such that it is hotter on one side than the other. Now it is intuitive to say that closer to the heat source, the temperature will be higher than further along the rod, where it will be colder. We can illustrate this temperature distribution with a function:

$$T(x, y, z)$$

where  $T$  is the temperature at a point in the rod, and  $(x,y,z)$  are the coordinates along the axis in 3 dimensions. This leads to these 3 partial derivatives:

$$\frac{\partial T}{\partial x}, \frac{\partial T}{\partial y}, \frac{\partial T}{\partial z}$$

These partial derivatives, gives us information about the direction and magnitude of heat flow at various points on the rod. The algorithm then searches for an equation  $T(x,y,z)$ , that sufficiently predicts the observed temperature distribution and it's partial derivatives, deriving laws such as the heat transfer equations, or elasticity relationships.

$$\frac{\partial T}{\partial t} = \alpha \nabla^2 T$$

Through using partial derivatives, we have in essence redefined the search criteria for the algorithm, through it's measure of the accuracy in comparison of potential solutions over the invariants represented in the experimental data [12] . This also leads to the pleasant finding, that it can additionally capture relationships that represent other identities of the system, beyond invariants and heat transfer equations.

You can subtly guide the type of laws that such an algorithm finds, by selectively picking the variables to input into the algorithm,. For example providing velocities and force to find laws of motion.

### 2.2.3 Dimensional Analysis:

Dimensional Analysis is a method of solving problems usually in maths and physics, where we analyse the relationships between different physical quantities, by comparing their "units." It is a powerful method of reducing the complexity of systems, enabling engineers and scientists to analyse problems that we can't even pose, much less solve the equations of [15].

Using the fact that numerous questions in science can be simplified by requiring the dimensions/units of the right and left hand side of the expression to be equal, we can transform the question into a smaller number of variables, which all have no dimension [7]. It has been automated to find the integer powers of expressions and has proven to be useful especially when the power is an irrational number.

Here is a general strategy that showcases how dimensional analysis can be used:

Let's say we have a variable in an equation that can be broken down into its fundamental units, such as (second, kilograms, ampere ...) to various powers. We can then take this, and represent each of the units as vectors, such that each of the fundamental units, is assigned a dimension, and it's important to note, this then allows us to represent any physical quantity as a product of these units, so let us construct a vector  $v$ , with 3 integers, where each corresponding integer represents the power of each of the fundamental units.

Given that we want to derive an expression, such as  $y = f(x_1, \dots, x_n)$  we can then create some matrix  $M$ . Each of the columns of the given matrix, is the unit vector  $v$  of the corresponding variable  $x_i$ . We then need to define another vector to represent the units of  $y$ , which will be called  $z$ . If we let the solution be some vector  $s$ , solving  $Ms = z$ , this then lets us raise the powers on both sides, to elevate the independent variables, to make this equation dimensionally consistent.

Taking the null space of the matrix  $M$ , where  $MV = 0$ , allows us a basis to create a dimensionless group, allows for a simplification of the problem.

This is also more intuitive to understand physical phenomena, the nature of physics comprehension, making this vital in further understanding derived laws, making the process easier to explain and understand [22]. Therefore, this is a crucial tool, for cultivating a deeper understanding of physics effectively [23].

### 2.2.4 Genetic Programming:

Genetic programming (GP), is a special evolutionary algorithmic technique, where the individuals are seen as programs that evolve, starting for a population, is iteratively "evolved," transforming the populations of individual programs, into other populations. This new generation of programs are created using some genetic operations or survival criteria, [8]mimicking natural evolutionary condition on earth.

A very basic overview, shows that genetic programming algorithms, consists of initializing the population [2], then evaluation of the said population through some predefined metrics and functions, followed by selection of the fittest programs based on the score given by the metric, and "genetic operation," such as reproduction, mutation and cross-over. The algorithm then iterates these steps thousands of times, through many generations, and finally terminates once the desired result has been achieved.

We can use genetic programming, and tweak the algorithm, and combine it with symbolic regression, to help us derive laws.

Consider modelling the various potential formulas as a tree, which is composed of various functions in the nodes. These functions can vary from arithmetic operations, mathematical functions, or defined unique opera-

tors. Then we can program the fitness function [5], and use it to measure how well the given potential expression in the population compares with the given databases, and given the nature of genetic programming, the better performing functions are more likely to be passed down into the next generation. Then after many iterations, we can give the solution with the best performance.

There are various ways to implement the fitness function, and for example we can use a criteria like this, along with mean squared error [14]:

$$V = 2X + N \cdot \ln(M/N)$$

Here M is the mean squared error, and N is the number of data points, X is the number of parameters used on the genetic programming algorithm. The lower the value of V is, the better the model performs. The performance of this strategy can then be evaluated with various other metrics, to judge how well the algorithm performs.

### 2.2.5 Broader Use Cases and Techniques:

SYMBBA [4] demonstrates using advanced Transformer models for complex symbolic computation in physics. It's relevant as an emerging alternative to traditional Symbolic Regression. Rather than searching expression space, it treats symbolic manipulation like language translation, directly mapping input sequences (problems) to output sequences (symbolic solutions) using powerful sequence-to-sequence techniques. This is essential for high-energy physics. Allowing physicists to rapidly generate theoretical predictions needed to interpret experimental data from particle colliders, speeding up scientific analysis and discovery.

## 3 PySR

This section describes the initial implementation of the symbolic regression models I developed using the PySR library.

### 3.1 Momentum Laws:

To generate the dataset, I chose 100 data points, and created two variables mass (M) and acceleration (a), each represented in two dimensions[9]. Then the data points were generated using *numpy.random.randn* function. The force (F), was then calculated to be the produce of these two data sets. Mass and acceleration were concatenated along the same axis using *numpy.concatenate*, resulting in a combined dataset. This is partially because the model used here, *PySRRegressor* expects a single array as input, and this helps highlight the relationship between these variables to the symbolic regression algorithm.

Then model performed symbolic regression, configured with 40 iterations along with a customer loss function, taken to be the squared difference between the prediction and the target variable.

$$\mathcal{L}(\hat{x}, x) = (\hat{x} - x)^2$$

The model was trained on this dataset, upon termination, it produced a list of potential candidate formulae, from which I manually identified the correct expression,  $F = m\ddot{a}$ .

---

**Algorithm 1:** Symbolic Regression for  $F = M \cdot A$ 

---

**Result:** A symbolic representation approximating  $F = M \cdot A$ **Initialization:**Generate random data for mass ( $M$ ) and acceleration ( $A$ );Compute target force values:  $F = M \cdot A$ ;Combine  $M$  and  $A$  into input matrix  $X$ ;**while** *Symbolic regression process* **do**

Train the symbolic regression model with the following settings;;

**Binary operators:** Multiplication (\*);    **Unary operators:** None;    **Loss function:** Mean squared error between predictions and targets;    **Iterations:** 40;    **if** *Current symbolic representation improves loss* **then**

Update the symbolic model;

Save the current best expression;

**else**

Continue exploration of new symbolic expressions;

**end****end**

---

Similarly, the other laws of momentum, were also derived using this approach.

$$\mathbf{F}\Delta t = \Delta \mathbf{p} = m(\mathbf{v}_f - \mathbf{v}_i) \quad (1)$$

$$m_1\mathbf{v}_{1,i} + m_2\mathbf{v}_{2,i} = m_1\mathbf{v}_{1,f} + m_2\mathbf{v}_{2,f} \quad (2)$$

### 3.2 Pendulum Laws:

The data is generated using numpy. The simulation involves, Euler's method to solve the pendulum's equation of motion. Through taking small and discrete steps, the method approximates the solution. The equation for a simple pendulum is given by:

$$\alpha = -\frac{g}{L} \sin(\theta) \quad (3)$$

 $\alpha$  angular acceleration ( $\text{rad/s}^2$ ) $g$  acceleration due to gravity ( $\text{m/s}^2$ ) $L$  length of the pendulum (m) $\theta$  angular displacement (rad)

The Euler technique approximates the changes in angular velocity and displacement over some small step in time, as follows:

$$\omega_{i+1} = \omega_i + \alpha_i \Delta t \quad (4)$$

$$\theta_{i+1} = \theta_i + \omega_{i+1} \Delta t \quad (5)$$

The function iterates through a few hundred time steps, updating the angular velocity and displacement at each time step. To prevent errors accumulating due to numerical drift, which are small errors that accumulate and become significant due to the inherent nature of approximation methods like Eulers. To keep the values coherent, a wrap around operation is used to ensure the angular displacement is within the range of  $[\pi, -\pi]$  radians.




Figure 3: The motion of a simple pendulum using the Euler method.

### 3.3 Noise:

This section investigates the model's robustness to noise. To simulate varying levels of interference, artificial noise was systematically introduced during the data generation phase, building upon the previously established model framework. Noise was modelled by generating random numbers within a range of progressively increasing magnitude, utilizing Python's random library. The objective was to observe and quantify the degradation in model accuracy as a function of increasing noise levels, as well as explore ways to mitigate it.

#### 3.3.1 How noise affects the model:

To introduce noise into the generated dataset, I imported Python's random library and used the randint function. To systematically vary the level of noise, I created an additional function that incrementally increased the parameters passed to randint, causing each successive dataset to become progressively noisier.

After generating these noisy datasets, the symbolic regression model was applied to each one, and the resulting equations were analyzed. The relationship between noise level and model performance was then visualized through a graph. Additionally, the time library was used to measure how long each run of the model took.

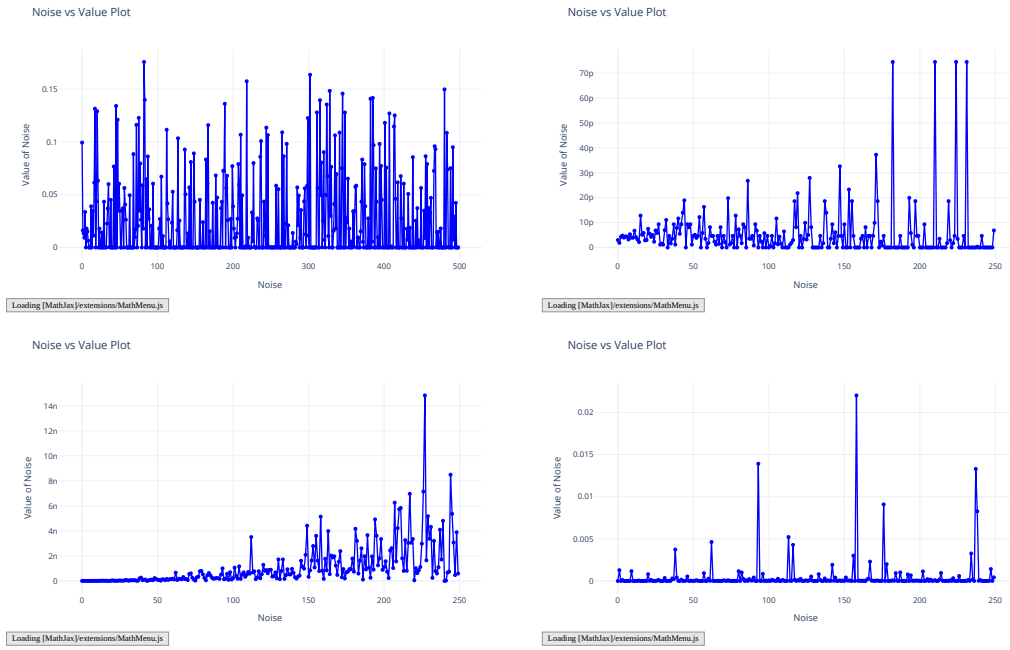


Figure 4: Plots showing how loss behaves as noise is increased randomly.

Shows from top left to bottom right, Conservation, Newtons first law, Newtons third law and Simple Harmonic Motion.

iterative:




Figure 5: Plots showing how loss behaves as noise is increased iteratively.

Shows from top left to bottom right, Conservation, Newtons first law, Newtons third law and Simple Harmonic Motion.

### 3.3.2 Denoise function:

Following the initial noise analysis, a subsequent experiment incorporated a denoising method (implemented using a Python library) applied to the data prior to processing by the pysr model. The results, as depicted in the accompanying plot, indicate that this denoising approach improved model performance up to a specific noise threshold. However, beyond this critical level, performance degraded comparably for both the denoised and non-denoised datasets, suggesting the denoise function's effectiveness diminished at higher noise intensities.[26]

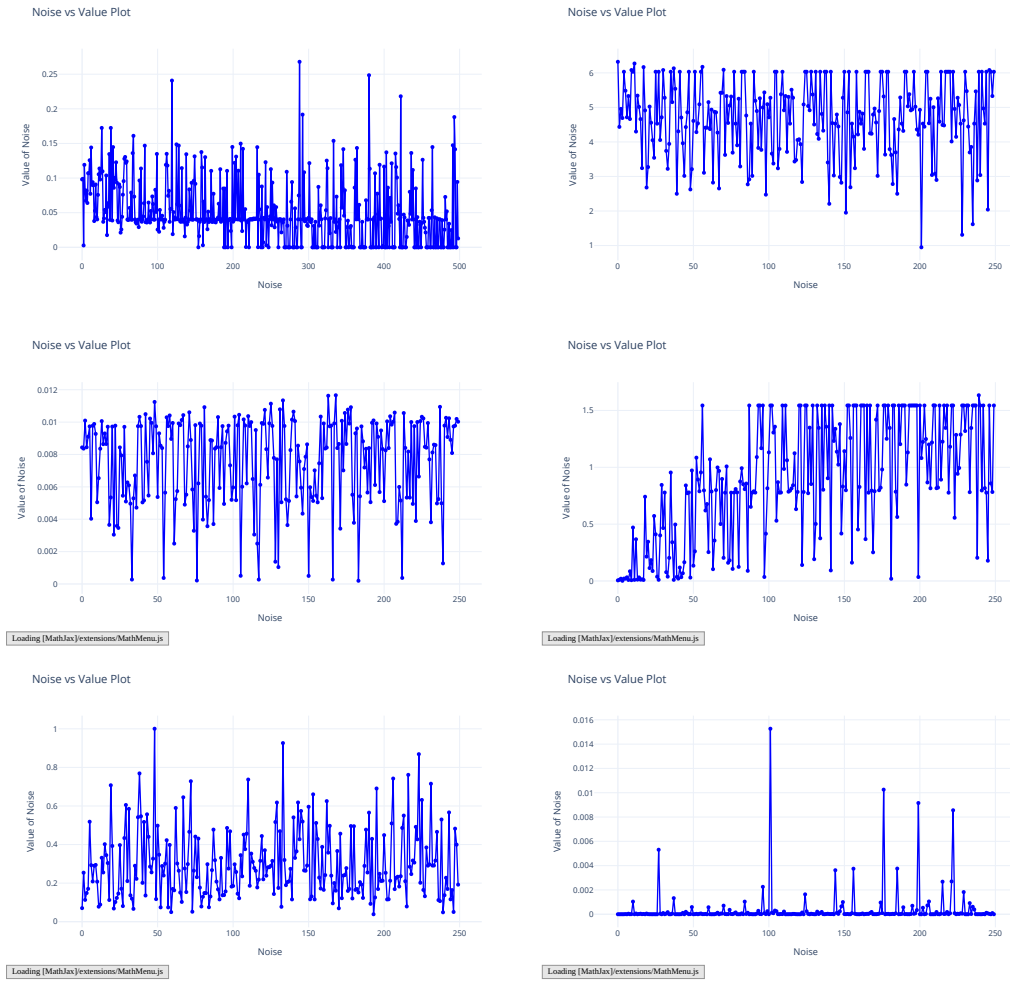


Figure 6: Noise plots with dimensionless reduction for various physical systems.

From top left to bottom right, shows Conservation, Newton's First, Newton's Second, Newton's Third, Rutherford, and Simple Harmonic Motion laws.

Iterative:






Figure 7: Noise plots with dimensionless reduction and iterative improvement for various physical systems.

From top left to bottom right, shows Conservation, Newton's First, Newton's Second, Newton's Third, Rutherford, and Simple Harmonic Motion laws.

## 4 Symbolic Regression from First Principles:

### 4.1 A Brute-Force Approach:

The core and essential component of any symbolic regression model lies in its ability to generate and traverse the search space of potential equations and expressions that best fit the given data. To streamline the process and validate the functionality of my expression generation, I began by implementing simple two-variable equations, specifying the operations used within the equation. This approach was initially limited to basic operations, with plans to extend it to accommodate constants and additional complexities.

---

**Algorithm 2:** Generate Initial Symbolic Expressions

---

**Data:** List of variable names, list of constants, list of operators

**Result:** List of generated symbolic expressions

**Initialization::**

Create pool of symbolic variables and constants from inputs;

Initialize empty list for expressions;

**for** each pair  $(a, b)$  from the pool (considering permutations) **do**

**for** each operator  $op$  from the list of operators **do**

        Form expression  $a \text{ op } b$ ;

**if** resulting expression is valid **then**

            Add expression to the list of expressions;

**end**

**end**

**end**

Return list of expressions;

---

Subsequently, I refined this approach by designing a recursive method to generate expressions, allowing for the creation of more robust and diverse equations from the available variables. This process is dynamic, making it adaptable to a variety of input configurations.

---

**Algorithm 3:** Recursive Generation of Symbolic Expressions

---

**Data:** List of operators, list of variable names, maximum depth

**Result:** List of symbolic expressions generated at the maximum depth

**Initialization::**

Create list 'Expressions[0]' containing symbolic variables from input;

**for** depth  $d$  from 1 to maximum depth **do**

    Create empty list 'Expressions[d]';

**for** each expression  $a$  in 'Expressions[d-1]' **do**

**for** each expression  $b$  in 'Expressions[d-1]' **do**

**for** each operator  $op$  in list of operators **do**

                Attempt to form new expression  $a \text{ op } b$ ;

**if** resulting expression is valid **then**

                    Add new expression to 'Expressions[d]';

**end**

**end**

**end**

**end**

**end**

Return 'Expressions[maximum depth]';

---

#### 4.1.1 Exploiting Physical Properties:

The next step involves truncating the generated expressions to prune the search tree as efficiently as possible. One effective method for achieving this is by leveraging the symmetrical properties of physical equations and recognizing their mathematical equivalence. This includes removing redundant or duplicate expressions that do not contribute new information.

This is the approach I used to achieve this:

#### 4.1.2 Dealing with constants:

Another approach I employed to further prune the set of generated expressions was by filtering out any expressions that did not contain all the specified variables. This step helps optimize the process by reducing the number of irrelevant expressions, ultimately saving computation time during the evaluation phase.

#### 4.1.3 Dealing with powers:

To apply powers to expressions, I incorporated the power operation directly into the generated expressions. This not only allows for the creation of more complex models but also facilitates further pruning of the search tree by avoiding the generation of redundant expressions that already incorporate powers.

---

**Algorithm 4:** Apply Powers Recursively to Expressions

---

**Data:** List of initial expressions, list of powers, maximum depth

**Result:** List of unique symbolic expressions generated by applying powers

**Initialization::**

Create set 'ResultExpressions' and add all initial expressions to it;

**for** depth  $d$  from 1 to maximum depth **do**

    Create empty set 'NewExpressionsThisDepth';

**for** each expression  $e$  in 'ResultExpressions' (from previous depths) **do**

**for** each power  $p$  in list of powers **do**

            Attempt to calculate  $\text{powered\_}e = e^p$ ;

**if** resulting expression is valid **then**

                Add  $\text{powered\_}e$  to 'NewExpressionsThisDepth';

**end**

**end**

**end**

    Add all expressions from 'NewExpressionsThisDepth' to 'ResultExpressions';

**end**

Return list representation of 'ResultExpressions';

---

I implemented a filtering mechanism based on whether the expression included a power operation. This further reduces the size of the search tree by eliminating unnecessary branches. While a more robust model would derive this power operation from scratch, I opted for this approach to optimize computation time and maintain flexibility.

---

**Algorithm 5:** Filter Expressions by Presence of Target Powers

---

**Data:** List of expressions, list of target powers

**Result:** List of expressions containing a sub-expression with a target power

**Initialization:**

Create empty list 'FilteredExpressions';

**for** each expression  $e$  in input expressions **do**

    Set 'found<sub>target power</sub>' to false *Traverse the symbolic tree of  $e$  (e.g., preorder);*

**for** each sub-expression  $s$  during traversal **do**

**if**  $s$  is a power AND exponent of  $s$  is in target powers **then**

            Set 'found<sub>target power</sub>' to true *Break traversal for  $e$ ;*

**end**

**end**

**if** 'found<sub>target power</sub>' is true **then**

        Add  $e$  to 'FilteredExpressions';

**end**

**end**

Return 'FilteredExpressions';

---

#### 4.1.4 Chaining powers and constants:

The next step was to chain together powers and constants, applying both to the generated expressions. While the existing design already supports chaining, it is essential to properly filter the results to ensure the search tree remains as compact as possible.

Although constants can be filtered using the existing method, the power operations are embedded within the constants, which causes the previous power filter to no longer function as expected. Consequently, I redesigned the filtering process to operate recursively, allowing it to handle both powers and constants effectively.

code:

However, this approach sometimes results in expressions that feature chained constants, such as  $\sin(\sin())$ . To refine the model further, I introduced an additional filter to remove expressions with multiple instances of the same constant chained together.

#### 4.1.5 Loading data:

Next, I needed an efficient way to load the data I had created. At this stage, my primary focus was on rapid testing. To facilitate this, I initially generated some dummy data values. Afterward, I decided to store the data as a NumPy array, as this would offer significant speed advantages over using text files. Several factors contribute to this, such as NumPy arrays being stored in memory, the efficiency of the underlying binary data format, and NumPy's use of C, which allows it to vectorize operations, greatly enhancing performance.

---

**Algorithm 6:** Load and Validate Data

---

**Data:** Input matrix  $X$ , target vector  $Y$ , list of variable names

**Result:** Input matrix  $X$ , target vector  $Y$ , list of symbolic variables

**Check Input Validity::**

**if** *number of columns in  $X$   $\neq$  number of variable names* **then**

    | Indicate error or stop execution;

**end**

**Process Variables::**

Create list of symbolic variables from the list of variable names;

Return  $X$ ,  $Y$ , and the list of symbolic variables;

---

As shown, I perform a check to ensure that the number of variables provided matches the shape of the array  $X$ , where  $X$  represents the input data, and  $y$  represents the target data, i.e., the final result. For example,  $X$  contains the mass and acceleration values, while  $y$  contains the corresponding values of the force  $f$  as calculated by the equation  $f=ma$ . This serves as a basic validation to confirm that the number of columns in the input data corresponds correctly to the variables provided.

#### 4.1.6 How to mitigate noise in data:

Ways to mitigate the noise and its affects on the model were explored. Functions such as "denoise," in the symbolic regression library helped to some extent. However after a certain point, such methods do not seem to offer much assistance.

I also made my own denoise algorithm. I implemented various different denoise algorithms to see what effects they had. Firstly I implemented a simple moving average as a way to mitigate the noise in the dataset. reword this -> "Simple and fast, smooths data well by averaging neighbors. However, it blurs sharp changes and is sensitive to extreme outlier values, pulling the average significantly and distorting the signal." These were my results, this is the pseudo code, explain the algorithm

The second denoise algorithm I implemented is a median filter, and this is what effects it has, and this is how i implemented it. Insert Pseudo code. reword: "Excellent at removing spikes and preserving edges better than averaging. Less affected by outliers. Can sometimes slightly distort the overall shape of the signal, especially with large window sizes."

Finally this is the third algorithm that I had implemented for denoising. Wavelet Denoising, this is the effects, and this is the pseudo code. Reword this -> "Transforms data to isolate noise, preserving both smooth and sharp signal features effectively. More complex to understand and requires careful selection of wavelet type and parameters for optimal results, which can be tricky."

#### 4.1.7 Evaluating expressions:

Next, I evaluate the expressions that have been generated. I assign the input variables to the corresponding columns of the data in increasing order. These values are then substituted into the expressions, and the model runs the calculations, producing an array of outputs for each expression. This process essentially evaluates every pruned expression and returns a NumPy array of results based on the input data.

---

**Algorithm 7:** Evaluate Symbolic Expression Numerically

---

**Data:** Symbolic expression, list of variable names, input data matrix  $X$

**Result:** Numerical evaluation of the expression for each data point in  $X$ , or NaN if evaluation fails

**Initialization::**

Create list of symbolic variables from the list of variable names;

**Convert to Numerical Function::**

Convert the symbolic expression into a numerical function, mapping symbolic variables to input columns of  $X$ ;

**Evaluate::**

Prepare input data from columns of  $X$  to match function arguments;

Evaluate the numerical function using the prepared input data;

Return the resulting array of values;

Evaluation fails (e.g., division by zero) Return an array of Not-a-Number (NaN) values with the same size as the number of data points;

---

Following the approach outlined in the paper [insert citation here], I utilized a medium error description length loss function, implementing it as described. The error is calculated using the squared difference to ensure all errors are positive, and a constant of 1 is added to guarantee that all errors are greater than 1 when taking the logarithm.

---

**Algorithm 8:** Calculate Mean Log Squared Error

---

**Data:** Array of predicted values, Array of original values

**Result:** Mean of the base-2 logarithm of (1 + squared error)

**Initialization::**

Flatten input arrays ‘predicted’ and ‘original’;

Initialize ‘total\_log\_error’ to 0.0 *Get the number of elements ‘n’ (length of arrays)*

**Calculate Total Log Error::**

**for** index  $i$  from 0 to  $n - 1$  **do**

    Calculate ‘error’ = absolute difference between ‘original[ $i$ ]’ and ‘predicted[ $i$ ]’;

    Calculate ‘squared\_error’ = ‘error’ squared Calculate ‘log\_error’ =

        base - 2 logarithm of (1 + ‘squared\_error’) Add ‘log\_error’ to ‘total\_log\_error’ **end**

**Calculate Mean Error::**

    Calculate ‘mean\_error’ = ‘total\_log\_error’ / ‘n’

    Return ‘mean\_error’

---

## 4.2 Polynomial Fit Module:

Now that the core of the algorithm is functional—handling constants, powers, variables, generating expressions, and filtering redundancy using physical principles like symmetry—I aimed to extend the program by implementing a polynomial fitting module. The goal of this technique is to efficiently fit data to a polynomial model, as many functions in physics (or parts of them) are well-approximated by low-order polynomials, and polynomial fitting is a computationally inexpensive method for this specific class of functions. The technique generates all possible polynomial terms up to a specified degree (e.g., degree 4) and creates a linear equation for each data point where the unknowns are the polynomial coefficients. The system of equations is solved using standard methods such as least squares, and the Root Mean Squared Error (RMSE) of the fit is calculated. If the RMSE is below a predefined tolerance (denoted as  $\text{pol}$ ), the polynomial is accepted as a solution. This approach serves as a fast base case in the recursive algorithm, quickly solving problems that are simple polynomials, and it can also handle sub-problems transformed into polynomial form by other modules, such as dimensional analysis or function inversion.

#### 4.2.1 Data Loading:

To begin, I developed the data loading function. The goal was to accept a NumPy array containing the data, along with a list of variables. The function then compares the shape of the data array with the number of variables provided to ensure that the input is consistent and sufficient for further processing.

---

**Algorithm 9:** Load and Validate Data Array

---

**Data:** Input data array, list of variable names

**Result:** Validated input data array (if valid)

**Check Input Validity::**

Get number of columns in data array;

Get number of variable names;

**if** *number of columns*  $\neq$  *number of variable names* **then**

    Raise a value error indicating the mismatch;

**end**

Return input data array;

---

#### 4.2.2 Generating polynomial expressions:

The next step involves generating polynomial expressions. The function returns a list of polynomial expressions based on the input coefficients, variables, and operators, considering a specified maximum degree for the terms.

The function works by first creating symbolic representations for the variables. It then iterates over all possible combinations of powers for the variables up to the specified degree and combines these terms using the provided operators. Finally, the generated expressions are simplified and returned as a list.

---

**Algorithm 10:** Generate Expressions from Terms and Operators

---

**Data:** List of coefficients, list of variable names, list of operators, maximum degree

**Result:** List of generated symbolic expressions

**Initialization::**

Create list of symbolic variables from variable names;

Initialize empty list 'GeneratedExpressions';

**Generate Expressions::**

**for** *each combination of powers (from 1 to max degree) for each variable* **do**

    Create a list of 'terms', where each term is 'coefficient \* variable

<sup>Power</sup> **for** *each combination of operators (one less than number of terms)* **do**

            Build an 'expression' by combining the 'terms' sequentially using the chosen operators;

            Simplify the 'expression';

            Add the simplified 'expression' to 'GeneratedExpressions';

**end**

**end**

Return 'GeneratedExpressions';

---

#### 4.2.3 Filtering the Polynomial expressions:

The filter\_expressions function programmatically filters symbolic expressions based on both structural and semantic constraints. It is particularly suited for large-scale symbolic filtering tasks where strict mathematical structures must be enforced.

The initial version of this function worked for symbolic constants (e.g., sin, cos, etc.) but failed to handle numbers or integer coefficients. This issue was identified during testing, prompting me to rewrite the function so

that it could also handle integer coefficients properly.

---

**Algorithm 11:** Filter Expressions by Variables, Constants, and Power

---

**Data:** List of expressions, List of required variable names, List of required constants (values/types),  
Required power value

**Result:** List of expressions matching all criteria

**Initialization::**

Create empty list 'FilteredExpressions';

Convert required variable names to symbolic variables;

**for** each expression  $e$  in input expressions **do**

    Set ' $vars_{ok}$ ' =  $true$  if all required symbolic variables are in  $e$ 's free symbols, false otherwise;

**if NOT** ' $vars_{ok}$ ' **then**

        | continue

**end**

    ;

    Set ' $constants_{ok}$ ' =  $true$  **for** each required constant  $c$  in list of required constants **do**

        Check if  $c$  is present as a sub-expression in  $e$  (matching value or type);

**if NOT** present **then**

            | ' $constants_{ok}$ ' =  $false$ ; break loop over constants **end**

            ;

**end**

**if NOT** ' $constants_{ok}$ ' **then**

            | continue

**end**

        ;

        Set ' $power_{ok}$ ' =  $false$  Check if any sub-expression in  $e$  is a power with exponent equal to  
        required power;

**if** present **then**

            | ' $power_{ok}$ ' =  $true$  **end**

            ;

**if NOT** ' $power_{ok}$ ' **then**

                | continue

**end**

            ;

            Add  $e$  to 'FilteredExpressions';

**end**

    Return 'FilteredExpressions';

---

#### 4.2.4 Evaluating expressions:

The next step involves fitting the filtered expressions to the dataset. The model fitting function fits polynomial expressions to the input data by determining the optimal set of coefficients that minimize the error between the predicted and actual output values. It evaluates multiple polynomial degrees, up to a specified maximum, and selects the degree that results in the lowest error, thereby ensuring an optimal balance between accuracy and complexity.

I utilize the Root Mean Squared Error (RMSE) to calculate the loss, and the function returns a list of loss values, one for each expression.



---

**Algorithm 12:** Evaluate Expressions and Calculate RMSE

---

**Data:** List of symbolic expressions, List of symbolic variables, Input data matrix  $X$ , True target vector  $Y_{true}$

**Result:** List of (expression, RMSE) pairs for successfully evaluated expressions

**Initialization:**

Create empty list 'EvaluationResults';

**for** each expression  $e$  in input expressions **do**

    Convert  $e$  into a numerical function using the symbolic variables;

    Evaluate the numerical function for each data point in  $X$  to get predicted values  $Y_{pred}$ ;

    Calculate Root Mean Squared Error (RMSE) between  $Y_{pred}$  and  $Y_{true}$ ;

    Add the pair ( $e$ , RMSE) to 'EvaluationResults';

    Evaluation fails (e.g., runtime error)

**end**

Return 'EvaluationResults';

---

To begin the fitting process, I take an expression, substitute the input variables with the corresponding values from the dataset, and compute the predicted  $y$ -values based on the equation. I then calculate the difference between the predicted values and the true target values ( $y$ ), which are the actual outputs. This difference is used to compute the Root Mean Squared Error (RMSE), which quantifies the prediction error for the expression.

#### 4.2.5 Best Polynomial Fit:

After calculating the RMSE values for each expression, I select the expression with the lowest RMSE as the most accurate polynomial fit for the data. This ensures that the chosen model has the best performance in terms of minimizing prediction error.

```
def bestFit(results): return min(results, key=lambda x: x[1])
```

### 4.3 Dimensional Analysis:

Physical equations must be dimensionally consistent, meaning the units on both sides of the equation must match, which severely limits the possible forms of the unknown function. This dimensional constraint provides a strong simplification of the problem, significantly narrowing the scope of valid equations. AI Feynman addresses this by applying dimensional analysis as the first step, simplifying the problem by identifying which combinations of variables are dimensionally consistent. The units of the variables—such as mass, length, and time—are represented as vectors of integer powers, forming a linear system based on the unit vectors of the input and target variables. Solving this system and finding the null space reveals dimensionless combinations of variables, which transforms the problem into one of finding a function of these new dimensionless variables. This process typically reduces the number of independent variables that the algorithm needs to search over, drastically shrinking the combinatorial search space for subsequent steps, such as polynomial fitting, brute force, and neural network-guided searches. As a result, the reduction in variables leads to a significant boost in efficiency, making these searches faster and more likely to succeed.

#### 4.3.1 Handling Units:

The AI Feynman database was accessed, and the units.csv file was downloaded to better understand the units present in the dataset. Upon reviewing the required units, a unit table was created in the form of an array, where each unit corresponds to a unique power of the fundamental SI units. Additionally, the basic SI units were implemented as an array/list to facilitate this mapping.

<i>SI Base Units</i>	<i>Vectorised Value</i>	<i>Derived Units</i>	<i>Vectorised Values</i>
<b>Mass</b>	[1, 0, 0, 0, 0, 0, 0]	<b>Area</b>	[0, 2, 0, 0, 0, 0, 0]
<b>Length</b>	[0, 1, 0, 0, 0, 0, 0]	<b>Acceleration</b>	[0, 1, -2, 0, 0, 0, 0]
<b>Time</b>	[0, 0, 1, 0, 0, 0, 0]	<b>Voltage</b>	[1, 2, -3, 0, -1, 0, 0]
<b>Temperature</b>	[0, 0, 0, 1, 0, 0, 0]	<b>Density</b>	[1, -3, 0, 0, 0, 0, 0]
<b>Current</b>	[0, 0, 0, 0, 1, 0, 0]	<b>Torque</b>	[1, 2, -2, 0, 0, 0, 0]
<b>Substance (mol)</b>	[0, 0, 0, 0, 0, 1, 0]	<b>Volume</b>	[0, 3, 0, 0, 0, 0, 0]
<b>Luminosity</b>	[0, 0, 0, 0, 0, 0, 1]	<b>Dimensionless</b>	[0, 0, 0, 0, 0, 0, 0]

Table 1: SI base and derived units with their corresponding vectorised representations.

There were also relevant derived units included.

#### 4.3.2 Construct Matrix and Target Vector:

This function constructs the dimensional matrix  $MM$  and the target vector  $bb$ , which are essential for performing dimensional analysis. It accepts lists of independent and dependent variable names, along with a dictionary that maps each variable name (as the key) to its corresponding unit vector (as the value). The unit vectors for the independent variables are retrieved through dictionary lookup, using lowercase variable names (i.e., `var.lower()`) to ensure case-insensitivity. These vectors are then efficiently assembled into the columns of matrix  $MM$  using `numpy.column_stack`, while the unit vector of the dependent variable forms the target vector  $bb$ . This approach ensures usability through case-insensitivity, leverages the performance benefits of `numpy.column_stack`, and includes explicit error handling to prevent issues arising from missing or incorrect keys.

---

#### Algorithm 13: Construct Dimensional Matrix and Target Vector

---

**Data:** List of independent variable names, Dependent variable name, Dictionary mapping variable names to dimensional vectors

**Result:** Matrix  $M$  of independent variable dimensional vectors, Vector  $b$  of dependent variable dimensional vector

**Construct Matrix  $M$ :**

Initialize empty matrix  $M$ ;

**for** each independent variable name  $v$  in the list **do**

    Look up the dimensional vector for  $v$  in the dictionary;

    Add this vector as a column to matrix  $M$ ;

    Variable name  $v$  not found in dictionary Raise an error indicating the missing independent variable;

**end**

**Construct Vector  $b$ :**

Look up the dimensional vector for the dependent variable name  $d$  in the dictionary;

Set this vector as vector  $b$ ;

Variable name  $d$  not found in dictionary Raise an error indicating the missing dependent variable;

Return matrix  $M$  and vector  $b$ ;

---

#### 4.3.3 Solving Dimension and Basis Units:

The `solveDimension` function solves the system of equations  $Mp=bMp=b$  for the unknown vector  $pp$ , where  $MM$  is the dimensional matrix and  $bb$  is the target vector. The function begins by converting the input matrices  $MM$  and  $bb$  into symbolic matrices using SymPy's `Matrix` class. It then attempts to solve for  $pp$  using the LU decomposition method (`LUsolve`), which is efficient for solving linear systems. If this process fails, an error

is raised, indicating the issue encountered during the solution attempt. Additionally, the function computes the null space of matrix  $MM$ , representing the set of dimensionless combinations of the variables. The function returns two outputs: the solution vector  $pp$  and the null space  $UU$ , which provides insight into any dimensionless combinations of the input variables. This method ensures robust error handling and leverages symbolic computation for accuracy.

---

**Algorithm 14:** Solve Dimensional System and Find Null Space

---

**Data:** Matrix  $M$ , Vector  $b$

**Result:** Particular solution vector  $p$ , Null space basis matrix  $U$

**Solve for Particular Solution  $p$ ::**

Solve the linear system  $M \cdot p = b$  for vector  $p$ ;

Solving fails (e.g., matrix is singular or system inconsistent) Raise an error indicating failure to solve the system;

**Calculate Null Space Basis  $U$ ::**

Calculate the basis vectors for the null space of matrix  $M$ ;

Form matrix  $U$  where columns are the null space basis vectors;

Return  $p$  and  $U$ ;

---

#### 4.3.4 Data Transformation Function:

The `generate_dimensionless_data` function is designed to transform a dataset into dimensionless form by applying the scaling factors derived from the solution vector  $pp$  and the null space  $UU$ . Initially, the function checks if  $pp$  is a NumPy array; if not, it converts it into a flattened NumPy array with `np.float64` type for consistency. The function then reshapes  $pp$  into a column vector and calculates a scaling factor by raising each variable in the input data `datax` to the powers specified by  $pp$ , followed by computing the product of these scaled values along the specified axis. This scaling factor is then used to adjust the target values `datay` to obtain the dimensionless target values `data_y_prime`.

If the null space  $UU$  is provided, the function proceeds to generate new dimensionless variables by computing the product of the input data `datax` raised to the powers specified by each vector in  $UU$ . These newly generated dimensionless variables are stacked together to form a transformed input dataset `data_x_prime`. If no null space is provided, the original input data `datax` is returned unchanged. This function provides a way to normalize and simplify the dataset, ensuring that all variables are dimensionless, which facilitates further analysis.

---

**Algorithm 15:** Generate Dimensionless Data

---

**Data:** Input data matrix  $X$ , Target data vector  $Y$ , Particular solution vector  $p$ , Null space basis matrix  $U$

**Result:** Dimensionless input matrix  $X'$ , Dimensionless target vector  $Y'$

**Prepare Particular Solution::**

Ensure  $p$  is a flattened numerical vector;

**Calculate Scaling Factor::**

Calculate a 'scaling\_factor' vector by raising each column of  $X$  to the corresponding power in  $p$  and taking the product across variables for each data point;

$$\text{scaling\_factor}_i = \prod_j X_{ij}^{p_j};$$

**Transform Target Data::**

Calculate dimensionless target vector  $Y'$  by dividing  $Y$  by the 'scaling\_factor'  $Y'_i = Y_i / \text{scaling\_factor}_i$ ;

**Transform Input Data::**

**if** Null space basis  $U$  is not empty **then**

    Create empty list 'DimensionlessVariables';

**for** each vector  $u$  in  $U$  **do**

        Calculate a new dimensionless variable vector by raising each column of  $X$  to the corresponding power in  $u$  and taking the product across variables for each data point;

$$\text{new\_var}_i = \prod_j X_{ij}^{u_j};$$

        Add 'new\_var' to 'DimensionlessVariables' **end**

        Form dimensionless input matrix  $X'$  by stacking the vectors in 'DimensionlessVariables';

**end**

**else**

        Set dimensionless input matrix  $X'$  equal to the original input matrix  $X$ ;

**end**

    Return  $X'$  and  $Y'$ ;

---

#### 4.3.5 Symbolic Transformation Generator:

This function generates the symbolic mathematical expressions corresponding to the dimensional analysis transformation. It accepts the original independent variable names (independent\_vars), the exponent vectors for scaling (p), and the dimensionless combinations (U).

The function first creates symbolic representations of each independent variable using sp.symbols from the SymPy library. Then, using the scaling exponents pp, it constructs the symbolic expression symbolic\_p representing the unit-fixing scaling factor xpxp through sp.Mul, which allows for the multiplication of terms. This expression effectively represents the scaling transformation applied to the input variables.

Next, the function iterates through each exponent vector uu in the null space UU, building the corresponding symbolic expressions for the dimensionless combinations. Each new variable is constructed by applying the powers from the exponent vector uu to the original input variables, and the resulting expressions are added to a list. This process ensures that both the scaling factors and the dimensionless combinations are represented as symbolic mathematical expressions, which are essential for understanding the relationship between the variables in the dimensional analysis.

---

**Algorithm 16:** Symbolic Transformation using Dimensional Analysis Results

---

**Data:** List of independent variable names, Particular solution vector  $p$ , Null space basis matrix  $U$

**Result:** Symbolic scaling factor expression, List of symbolic dimensionless group expressions

**Initialization::**

Create list of symbolic variables from independent variable names;

**Construct Symbolic Scaling Factor::**

Form symbolic expression for scaling factor by taking the product of each symbolic variable raised to the corresponding power in  $p$ ;

ScalingFactor =  $\prod_i \text{variable}_i^{p_i}$ ;

**Construct Symbolic Dimensionless Groups::**

Initialize empty list 'SymbolicDimensionlessGroups';

**for** each vector  $u$  in Null space basis  $U$  **do**

    Form symbolic expression for a dimensionless group by taking the product of each symbolic variable raised to the corresponding power in  $u$ ;

    DimensionlessGroup $_u$  =  $\prod_i \text{variable}_i^{u_i}$ ;

    Add 'DimensionlessGroup $_u$ ' to 'SymbolicDimensionlessGroups';

**end**

Return 'ScalingFactor' and 'SymbolicDimensionlessGroups';

---

## 4.4 Neural Network Fitting:

The next critical component involves using neural networks to predict and compute gradients from the dimensionless data, providing valuable insights for visualization. While neural networks do not directly solve for symbolic expressions, they serve as powerful tools for approximating complex relationships within the data. By training a neural network on the dimensionless variables, it can predict the output for a given input and calculate gradients, offering a smooth, differentiable function that helps visualize how changes in the input variables influence the model's predictions [4]. This capability allows us to gain a deeper understanding of the underlying behavior of the system. Although neural networks do not offer an explicit symbolic expression, they provide a flexible and efficient way to visualize the functional dependencies, aiding in the interpretation of complex patterns that may be difficult to express symbolically. Ultimately, the network's predictions and gradients can be used to explore and understand the data, even if the true functional form remains implicit.

### 4.4.1 Symbolic Network:

This class defines the neural network architecture used as a universal function approximator within the symbolic regression framework. It inherits from `torch.nn.Module`, the base class for all neural network modules in PyTorch. The constructor (`init`) initializes the network structure, accepting the number of input features (`n_input`) and defaulting to a single output (`n_output`). The model consists of several fully connected layers, with each layer followed by a Tanh activation function, which was selected based on the specifications of the reference papers. This multi-layer perceptron (MLP) architecture provides significant expressive power [1], making it capable of approximating complex relationships within the data. While this project approaches this by being used for dimensionless data, you can also use neural networks as sequential weak learners in gradient boosting frameworks to build approximations [6]. The Tanh activation ensures smooth, differentiable non-linearity, making the network suitable for approximating complex physical functions. While the neural network does not directly solve symbolic expressions, it enables us to predict outputs and compute gradients from the dimensionless data. These predictions and gradients help visualize how changes in input variables affect the model's outputs, offering valuable insights into the system's underlying behavior and functional relationships. This approach aids in interpreting complex patterns, even when the explicit symbolic form is not easily obtainable.

---

**Algorithm 17:** Symbolic Regression Neural Network Architecture

---

**Class Definition::**

Define a Neural Network Class ‘SymbolicNetwork’ inheriting from a base Neural Network Module;

**Initialization Method** ( *$n_{init}$* , *Number of input features*(‘ $n_{input}$ ’), *Number of output features*(‘ $n_{output}$ ’, default 1)

;

Call the constructor of the base Neural Network Module;

Define a sequential model containing the following layers::

Linear layer: ‘ $n_{input}$ ’ inputs, 128 outputs Tanh activation function Linear layer :

128 inputs, 128 outputs Tanh activation function Linear layer :

128 inputs, 64 outputs Tanh activation function Linear layer :

64 inputs, 64 outputs Tanh activation function Linear layer : 64 inputs, ‘ $n_{output}$ ’ outputs

**Forward Method** (‘forward’):;

**Data:** Input tensor ‘x’

;

**Result:** Output tensor

;

Pass the input tensor ‘x’ through the defined sequential model;

Return the output tensor;

---

#### 4.4.2 Preparing the data:

This function preprocesses raw input (data\_x) and output (data\_y) data into a format suitable for PyTorch model training and validation. It begins by converting the input NumPy arrays into PyTorch tensors of type torch.float32, as PyTorch requires tensor structures for efficient computation. The data is then split into training and validation sets based on a configurable train\_split ratio, ensuring the model can be evaluated on unseen data to prevent overfitting. The TensorDataset class is used to pair the input and target data, enabling efficient indexing during training, while the DataLoader handles batching, memory management, and parallel processing. Additionally, the DataLoader automates data shuffling, which improves model robustness by ensuring the model does not memorize the data order, aiding in better convergence during training. This design follows standard PyTorch practices to ensure compatibility, scalability, and efficient handling of large datasets, while the training-validation split is crucial for monitoring generalization and model performance.

---

**Algorithm 18:** Prepare Data for Neural Network Training

---

**Data:** Input data array  $X$ , Target data array  $Y$ , Batch size, Training split ratio (default 0.8)

**Result:** Training data loader, Validation data loader

**Convert to Tensors::**

Convert  $X$  to a PyTorch tensor (float32);

Convert  $Y$  to a PyTorch tensor (float32) and add a dimension;

**Split Data::**

Get the total number of samples;

Calculate the index for splitting based on the training split ratio;

Split the input and target tensors into training and validation sets;

**Create Datasets::**

Create a training dataset from the training tensors;

Create a validation dataset from the validation tensors;

**Create Data Loaders::**

Create a training data loader from the training dataset, using the specified batch size and enabling shuffling;

Create a validation data loader from the validation dataset, using the specified batch size and disabling shuffling;

Return the training data loader and the validation data loader;

---

#### 4.4.3 Training the Network:

This function orchestrates the supervised training process for the provided PyTorch neural network model, with the primary goal of adjusting the model's parameters (weights and biases) to minimize the difference between its predictions and the true target values using the training data, while also monitoring performance on unseen validation data. The function starts by transferring the model to the specified compute device (either 'cpu' or 'cuda') and initializes the Adam optimizer, a commonly used adaptive learning rate optimization algorithm. The optimizer is linked to the model's parameters, with the learning rate set as a hyperparameter. The Mean Squared Error (MSE) loss function is used, which is well-suited for regression tasks as it minimizes the squared error between predicted and actual values. The training loop iterates through multiple epochs, where the model is trained using batches from the training data, and the optimizer updates the model's parameters based on the gradients computed from the loss function. During the validation phase, the model's performance is evaluated without gradient calculation to save computational resources. This standard PyTorch training loop, which uses DataLoader for efficient batch processing, ensures the model is correctly trained and evaluated on both training and validation datasets, with printed outputs for tracking progress.

---

**Algorithm 19:** Train Neural Network Model

---

**Data:** Neural network model, Training data loader, Validation data loader, Number of epochs, Learning rate, Device

**Result:** Trained neural network model (implicitly returned by modification)

**Initialization::**

Move model to the specified device;

Initialize an optimizer (e.g., Adam) with model parameters and learning rate;

Define a loss function (e.g., Mean Squared Error);

**for** *each epoch from 1 to Number of epochs* **do**

**Training Phase::**

        Set model to training mode;

        Initialize total training loss for the epoch;

**for** *each batch (inputs, targets) in Training data loader* **do**

            Move inputs and targets to the specified device;

            Perform forward pass to get predictions;

            Calculate loss between predictions and targets;

            Zero gradients of model parameters;

            Perform backward pass to compute gradients;

            Update model parameters using the optimizer;

            Accumulate batch loss to total training loss;

**end**

        Calculate average training loss for the epoch;

**Validation Phase::**

            Set model to evaluation mode;

            Initialize total validation loss for the epoch;

            Disable gradient calculations;

**for** *each batch (inputs, targets) in Validation data loader* **do**

                Move inputs and targets to the specified device;

                Perform forward pass to get predictions;

                Calculate loss between predictions and targets;

                Accumulate batch loss to total validation loss;

**end**

            Calculate average validation loss for the epoch;

            Enable gradient calculations;

            Print epoch number, training loss, and validation loss;

**end**

---

#### 4.4.4 Predict Function:

This function performs inference using a trained PyTorch model, generating output predictions for a given set of input data. It is designed to take a trained model, input data as a NumPy array (`x_numpy`), and the target computation device (`'cpu'` or `'cuda'`). The function begins by setting the model to evaluation mode using `model.eval()`, which is critical as it disables layers like dropout or batch normalization that behave differently during training and inference, ensuring deterministic output. The input data (`x_numpy`) is then converted into a PyTorch tensor with the appropriate `dtype=torch.float32` and transferred to the specified device. The core prediction step occurs within a `with torch.no_grad()` context manager, which disables gradient calculation, thereby reducing memory consumption and speeding up computations during inference. This design follows standard PyTorch inference practices, ensuring that the model behaves correctly during prediction and optimizes performance by avoiding unnecessary gradient computations. Finally, the predictions are returned as a NumPy array, providing a user-friendly output format that is compatible with typical post-processing and analysis workflows.



---

**Algorithm 20:** Generate Predictions using Trained Network

---

**Data:** Trained neural network model, Input data (NumPy array), Device

**Result:** Predicted output (NumPy array)

**Setup::**

Set model to evaluation mode;

Convert input data (NumPy array) to PyTorch tensor (float32) and move to the specified device;

**Predict::**

Disable gradient calculations;

Perform forward pass using the model with the input tensor;

Move the resulting predictions tensor to CPU;

Convert the predictions tensor to a NumPy array;

Enable gradient calculations;

Return the predicted NumPy array;

---

## 4.5 Pareto Frontier Optimisation:

The Pareto frontier provides a valuable method for balancing the trade-off between accuracy and simplicity (complexity) in symbolic regression. Rather than searching for a single "best" formula, the objective is to identify all Pareto-optimal formulas, which are those for which no other formula is both simpler and more accurate. To implement this, a 2D plot is created where the x-axis represents the formula's complexity and the y-axis represents the mean error description length. As candidate formulas are generated through methods such as brute-force search, transformations, or recursive combinations, each candidate is evaluated and plotted as a point on the frontier. A critical aspect of this approach is Pareto pruning: any candidate formula that is dominated – meaning that another formula on the frontier has both equal or lower complexity and equal or lower loss (with at least one strict inequality) – is discarded just like in [9]. This pruning process reduces the search space and improves computational efficiency by eliminating redundant solutions, especially when combining results from subproblems. Additionally, it enhances robustness against noise by inherently favoring simpler formulas that achieve a given accuracy. The final Pareto frontier provides a spectrum of optimal solutions, offering the user a range of trade-offs between complexity and accuracy [8].

### 4.5.1 Points:

The Point class is a fundamental data structure within the Pareto frontier analysis framework, encapsulating the key characteristics of a single candidate formula evaluated during the symbolic regression process. Its constructor initializes each instance with three attributes: complexity, a numerical score representing the formula's complexity; accuracy, a measure of the formula's performance; and the formula itself, which is the associated symbolic expression. This design consolidates the essential metrics (complexity, accuracy) needed for Pareto dominance checks into a single object, simplifying the management and comparison of candidate solutions. Functions operating on the Pareto frontier can efficiently access the complexity and accuracy attributes to determine dominance relationships, facilitating the maintenance of a non-dominated set of optimal formulas.

---

**Algorithm 21:** Point Class for Pareto Front

---

**Class Definition;**

Define a Class ‘Point’ to represent a candidate solution;

**Initialization Method** (*init*,): *Complexityvalue, Accuracyvalue, Symbolicexpression*

;

Store the input complexity, accuracy, and expression as attributes of the Point object;

**Representation Method** (*repr*,): *StringrepresentationofthePointobject*

;

Construct a descriptive string including the complexity, accuracy, and expression values;

Return the constructed string;

---

**4.5.2 Pareto Point Set:**

The `update_pareto_points` function maintains a Pareto frontier of symbolic expressions, balancing model complexity and accuracy. It filters a combined list of candidate solutions to retain only non-dominated points, where a point A dominates point B if A is both no more complex and no less accurate, with at least one of these metrics strictly better. This approach ensures that only the best trade-offs between complexity and accuracy remain. In symbolic regression, this is crucial as the goal is to identify expressions that are both accurate and interpretable—low in complexity. The Pareto frontier provides a valuable tool for visualizing and selecting optimal models, helping to avoid overfitting (excessive complexity) and underfitting (excessive simplicity).

---

**Algorithm 22:** Update Pareto Front with New Candidate Solutions

---

**Data:** List of current Pareto points (complexity, accuracy, expression), List of new expressions, List of corresponding losses

**Result:** List of unique Pareto points (complexity, accuracy, expression) from the combined set

**Prepare New Points::**

Create empty list 'NewPoints';

**for** each new expression  $e$  and its loss  $l$  **do**

    Calculate complexity of  $e$ ;

    Add (complexity,  $l$ ,  $e$ ) as a tuple to 'NewPoints';

**end**

**Combine and Identify Pareto Front::**

Combine 'CurrentPoints' and 'NewPoints' into a single list 'CombinedPoints';

Create empty list 'ParetoPoints';

**for** each point  $p_1$  in 'CombinedPoints' **do**

    Set 'is\_dominated' to false **for** each other point  $p_2$  in 'CombinedPoints' ( $p_2 \neq p_1$ ) **do**

**if**  $p_2$  dominates  $p_1$  (better or equal in all objectives, strictly better in at least one) **then**

            Set 'is\_dominated' to true **break inner loop** **end**

**end**

**if NOT** 'is\_dominated' **then**

            Add  $p_1$  to 'ParetoPoints';

**end**

**end**

**Remove Duplicates::**

Create empty set 'SeenPointsKeys';

Create empty list 'UniqueParetoPoints';

**for** each point  $p$  in 'ParetoPoints' **do**

    Create a unique key for  $p$  (e.g., based on complexity, loss, and string representation of expression);

**if** key is NOT in 'SeenPointsKeys' **then**

        Add  $p$  to 'UniqueParetoPoints';

        Add key to 'SeenPointsKeys';

**end**

**end**

Return 'UniqueParetoPoints';

---

## 4.6 Plotting:

The need arose to find ways to visualize and better understand the processes occurring behind the scenes in the symbolic regression program. Humans are inherently visual creatures, and visualizations provide a more powerful means of comprehending data compared to simply examining strings of numbers or expressions, which makes it difficult to identify patterns or errors manually. Plotly was chosen over the more commonly used Matplotlib for its superior aesthetic appeal and its enhanced interactive environment, offering a more intuitive and engaging experience for exploring the data and model results.

### 4.6.1 Plot for Pareto front:

Visualizing the Pareto frontier is crucial for interpreting the results of symbolic regression as it offers an intuitive graphical representation of the trade-off between model complexity and predictive accuracy (or loss). The provided code leverages the Plotly Express library to generate this visualization. A Pandas DataFrame (df) is used to organize the data, containing columns for 'Complexity' (x-axis), 'Loss' (y-axis, representing inaccuracies such as MEDL), and the corresponding symbolic 'Expression'. The px.scatter function generates a scatter

plot that maps complexity against loss, with the 'Expression' data being displayed as labels on each point using the text argument. Additional customization via fig.update\_traces enhances the visibility of the points (marker size, color) and optimizes label positioning. This visualization allows researchers to easily identify the set of non-dominated solutions. By plotting quantitative metrics, it reveals where significant accuracy improvements require substantial increases in complexity (often visible in the "elbow" region), aiding in model selection based on the desired balance between interpretability and predictive power. Direct labeling of points with their expressions provides immediate context for each optimal solution.

---

**Algorithm 23:** Plot Pareto Frontier using Plotly

---

**Data:** List of Pareto points (complexity, accuracy, expression), Plot title (default)

**Result:** Plotly figure object representing the Pareto frontier

**Prepare Data::**

Convert the list of Pareto points into a structured data format (e.g., a table or DataFrame) with columns for Complexity, Loss, and Expression;

**Create Scatter Plot::**

Generate a scatter plot using the prepared data;

Map Complexity to the x-axis, Loss to the y-axis;

Use Expression as text labels for each point;

Set the plot title;

**Style Plot::**

Customize the appearance of the data points (e.g., size, color);

Customize the appearance of the text labels (e.g., position);

Customize the plot layout (e.g., background colors, font, title alignment);

Return the generated Plotly figure object;

---

## 4.7 Main method bringing it together Regressor:

### 4.7.1 main solver

#### Part 1: Dimensional Analysis and Symbolic Scaling

The function first applies Dimensional Analysis (DA) to reduce the problem's dimensionality. It computes a transformation matrix and scaling vector by solving the unit balance equation between independent and target variables. If no dimensionless groups are found, the result is a simple scaling law expressed symbolically. Otherwise, it computes and displays the dimensionless groups and the scaling part of the solution. This ensures the model operates on physically meaningful, unit-consistent quantities, simplifying the functional search space and preventing non-physical relationships. If DA completely solves the problem, the function terminates here by outputting the symbolic solution [24].

#### Part 2: Neural Network Approximation of Dimensionless Relation

When dimensionless groups exist, the function generates dimensionless data and fits a lightweight symbolic neural network to model the relationship between transformed inputs and outputs. A standard training pipeline is executed: data preparation, model instantiation, training with backpropagation, and prediction. The model's mean squared error (MSE) is computed on the dimensionless targets to evaluate fit quality. Neural network outputs (predictions and gradients) serve as a flexible preliminary approximation, potentially capturing nonlinear relations that guide the subsequent, more rigid symbolic regression steps.

#### Part 3: Polynomial Candidate Generation and Filtering

Following the neural fit, the function applies Polynomial Fitting (PF) techniques to generate candidate symbolic expressions [24]. Polynomial terms are systematically composed using provided variables, coefficients, and operators up to a specified degree. Generated expressions are filtered based on structural criteria, such as variable presence and coefficient validity, to ensure physical plausibility. Each polynomial is evaluated against the original data, and a best-fit candidate is selected according to its error score. If a candidate perfectly matches (zero error), the function outputs the discovered symbolic expression and terminates.

#### Part 4: Brute Force Symbolic Search and Evaluation

If no polynomial yields a perfect fit, the function initiates a Brute Force (BF) symbolic search. It exhaustively generates expressions by combining variables, constants, operators, and powers. Multiple filters (symmetry, variable relevance, power range, constant handling) systematically prune the expression set to reduce computational complexity and preserve physically meaningful candidates. The remaining expressions are evaluated against the dataset, seeking the best match based on performance metrics. This final exhaustive step ensures that even highly non-obvious symbolic relations can be discovered if they exist within the defined operator and degree space.

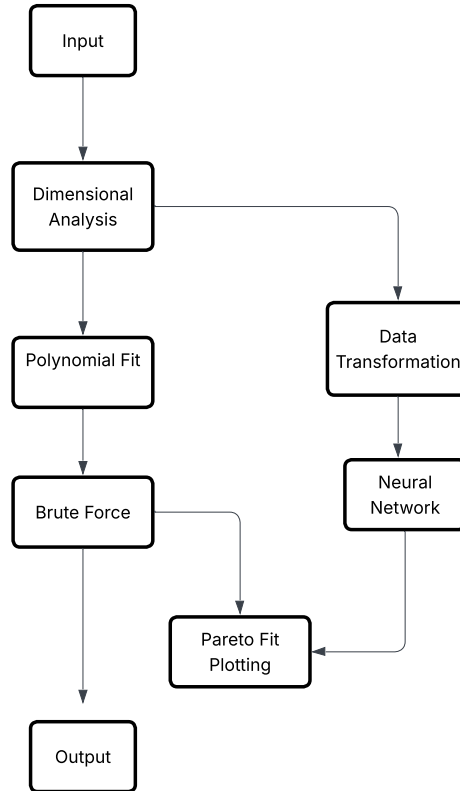


Figure 8: The SR control flow for the solver, show in this diagram above.

## 4.8 Testing:

Every module, function and file, were thoroughly tested using dedicated tests. The boundary conditions and inserted tests to make sure the function behaved as envisioned. There was also significant robust testing functions written when chaining together various techniques and models, in order to ensure everything was working smoothly.

### 4.8.1 Unit Testing:

I've written unit tests for every function, every evaluation method and have rigorously tested their functionality, edge cases, normal inputs etc.

<b>Module:</b>	<b>Tests:</b>	<b>Results:</b>
Dimensional Analysis	Get, Solve, Generate, Symbolic Transformation	Passed
Neural Network	Load, Network, Train, Predict, Gradient	Passed
Plotting	Pareto, Gradient	Passed
Polynomial Regressor	Load, Generate, Filter, Evaluate, Best fit	Passed
Brute Force	Load, Generate, Recursive, Evaluate, Loss, Variable Filtering, Constants, Powers	Passed
Main Solver	Main Solver, Laws	Passed

Rigorous unit testing was employed to validate the correctness of each module within the symbolic regression framework. Key components including Pareto frontier operations, polynomial generation and evaluation, and dimensional analysis routines were individually tested using structured unit tests. Each test assessed expected outputs under standard and edge-case inputs.

Test results confirmed correct behavior, and outputs aligned with theoretical expectations (e.g., correct Pareto set retrieval, successful polynomial fitting on synthetic data, and dimensionally consistent variable transformations). A summary of testing outcomes is shown below. Furthermore, synthetic datasets based on known biological formulas were used to validate the regressor’s ability to recover exact symbolic relationships, demonstrating the system’s practical reliability.

#### 4.8.2 Integration Testing:

To validate the interoperability and functional correctness of the developed software system, a comprehensive suite of integration tests was designed and executed. These tests focused on the interactions between the core Python modules, including Dimensional Analysis (`dimensionalAnalysis`), Polynomial Fitting (`polynomialFit`), Brute-Force Symbolic Regression (`bruteForce`), Neural Network (`neuralNetwork`) components, Pareto front optimization (`pareto`), and visualization (`plots`). Utilizing Python’s `unittest` framework, the tests simulated realistic data analysis workflows, verifying the seamless flow of data and control between modules. This type of chaining multiple methods was partly inspired by [9], showing that it is far more effective to employ this architecture when aiming to prune the search tree.

Key tested interactions included the preprocessing of data using Dimensional Analysis feeding into various model fitting approaches (PF, BF, NN), the evaluation of candidate expressions generated by BF and PF modules and their subsequent management on a Pareto front, and the symbolic transformation capabilities of the Dimensional Analysis module. Further tests confirmed the internal consistency of complex operations within the Brute-Force module and the functionality of gradient calculations in the Neural Network. Successful execution of these tests demonstrates the robustness of the integrated system and its capability to perform complex, multi-stage symbolic regression tasks as designed, ensuring reliable data handoffs and component communication within the software architecture.

#### 4.8.3 Performance Testing:

To quantitatively evaluate the computational efficiency and scalability of the core symbolic regression components, dedicated performance tests were executed on the `bruteForce` and `polynomialFit` modules. Utilizing Python scripts leveraging the `time` module for execution duration measurement and the `psutil` library for monitoring Resident Set Size (RSS) memory usage, these tests targeted the primary computational bottlenecks: symbolic expression generation and expression evaluation against numerical data.

The methodology involved systematically varying key parameters influencing complexity, such as expression generation depth (`bruteForce`), maximum polynomial degree (`polynomialFit`), the number of input variables, and the size of the input data sample for evaluation. By recording execution time and memory consumption

under these varying conditions, the tests aimed to characterize the performance scaling of each module’s algorithms. This analysis provides crucial insights into the computational complexity and resource requirements associated with each approach, enabling an understanding of their practical limitations and scalability characteristics. The results quantify the performance trade-offs inherent in different symbolic search strategies and inform the feasibility of applying the developed system to large-scale scientific discovery tasks by establishing empirical benchmarks for time and memory demands.

## 5 Applying the model to Biological Data:

However, its application to biological systems remains an underexplored frontier. Biological processes are inherently noisy, high-dimensional, and complex, often lacking explicit governing equations. As a result, there exists a significant research gap in employing symbolic regression techniques [25] for biological data modeling. A limited number of studies attempt this integration, making it a novel and potentially revolutionary field of study.

In this work, we apply symbolic regression to a fundamental biological process: the prediction of nucleic acid melting temperature ( $T_m$ ) from DNA/RNA base counts. The melting temperature is computed through a well-established empirical formula known as the Wallace rule, defined as:

$$T_m = 2(A+T) + 4(G+C)$$

where A,T,G,C denote the counts of adenine, thymine, guanine, and cytosine bases, respectively.

To test the approach, synthetic data is generated programmatically. Arrays are constructed where each sample consists of four integer inputs corresponding to base counts. The target  $T_m$  values are computed exactly according to the Wallace formula. Variables (`'A'`, `'T'`, `'G'`, `'C'`) and constants (`[2, 4]`) are explicitly defined to align with the biological context.

The symbolic regression system first performs dimensional analysis, even though biological units are less rigorously defined compared to physics. Then, using a neural-symbolic approach followed by polynomial and brute-force search stages, the system attempts to recover the original expression.

This study highlights how biological systems can be rendered into symbolic, mechanistic expressions, bridging the gap between empirical observation and interpretable models, and opening new avenues for data-driven biological discovery.

## 6 Evaluation:

## References

- [1] Ismail Alaoui Abdellaoui and Siamak Mehrkanoon. Symbolic regression for scientific discovery: an application to wind speed forecasting, 2021.
- [2] Mahed Abroshan, Saumitra Mishra, and Mohammad Mahdi Khalili. Symbolic metamodels for interpreting black-boxes using primitive functions, 2023.
- [3] Guilherme Seidyo Imai Aldeia and Fabrício Olivetti de França. Interpretability in symbolic regression: a benchmark of explanatory methods using the feynman data set. *Genetic Programming and Evolvable Machines*, 23, 2022.
- [4] Abdulhakim Alnuqaydan, Sergei Gleyzer, and Harrison Prosper. Symba: symbolic computation of squared amplitudes in high energy physics with machine learning. *Machine Learning: Science and Technology*, 4(1):015007, January 2023.
- [5] Peter J. Angeline. Genetic programming: On the programming of computers by means of natural selection,. *Biosystems*, 33, 1994.
- [6] Sarkhan Badirli, Xuanqing Liu, Zhengming Xing, Avradeep Bhowmik, Khoa Doan, and Sathiya S. Keerthi. Gradient boosting neural networks: Grownnet, 2020.
- [7] W. Blasiak, M. Godlewska, R. Rosiek, and D. Wcislo. Spectrum of physics comprehension. *European Journal of Physics*, 33, 2012.
- [8] William La Cava, Patryk Orzechowski, Bogdan Burlacu, Fabrício Olivetti de França, Marco Virgolin, Ying Jin, Michael Kommenda, and Jason H. Moore. Contemporary symbolic regression methods and their relative performance, 2021.
- [9] Miles Cranmer. Interpretable machine learning for science with pysr and symbolicregression.jl, 2023.
- [10] Miles Cranmer, Alvaro Sanchez-Gonzalez, Peter Battaglia, Rui Xu, Kyle Cranmer, David Spergel, and Shirley Ho. Discovering symbolic models from deep learning with inductive biases. In *Advances in Neural Information Processing Systems*, volume 2020-December, 2020.
- [11] Tom Hope, Doug Downey, Daniel S. Weld, Oren Etzioni, and Eric Horvitz. A computational inflection for scientific discovery. *Communications of the ACM*, 66, 2023.
- [12] G. V. Kelly. Partial differential equations for scientists and engineers by g. stephenson. *Irish Mathematical Society Bulletin*, 0018, 2021.
- [13] Alexandre Koyré and R. E.W. Maddison. *The astronomical revolution: Copernicus - Kepler - Borelli*, volume 20. online, 2013.
- [14] Andrew R. Liddle, Pia Mukherjee, and David Parkinson. *Model selection and multi-model inference*, volume 9780521887946, chapter 2. online, 2009.
- [15] Sandro G. Longo. *Dimensional Analysis*, volume 05, chapter 5. online, 2021.
- [16] Nour Makke and Sanjay Chawla. Interpretable scientific discovery with symbolic regression: a review. *Artificial Intelligence Review*, 57, 2024.
- [17] Daniel Rivero, Enrique Fernandez-Blanco, and Alejandro Pazos. Dome: A deterministic technique for equation development and symbolic regression. *Expert Systems with Applications*, 198, 2022.
- [18] Michael Schmidt and Hod Lipson. Distilling free-form natural laws from experimental data. *Science*, 324, 2009.



- [19] Michael Schmidt and Hod Lipson. Supporting online material for distilling free-form natural laws from experimental data. *Science (New York, N.Y.)*, 324, 2009.
- [20] R.T. Smith. Calculus early transcendentals functions. *Calculus Early Transcendentals Functions 4TH edition*, 53, 2012.
- [21] James Stewart. *Essential Calculus: Early Transcendentals*. online, 2012.
- [22] Keith S. Taber. Maths should be the last thing we teach. *Physics Education*, 44, 2009.
- [23] Wassim Tenachi, Rodrigo Ibata, and Foivos I. Diakogiannis. Deep symbolic regression for physics guided by units constraints: Toward the automated discovery of physical laws. *The Astrophysical Journal*, 959, 2023.
- [24] Silviu-Marian Udrescu and Max Tegmark. Ai feynman: a physics-inspired method for symbolic regression, 2020.
- [25] Jared Willard, Xiaowei Jia, Shaoming Xu, Michael Steinbach, and Vipin Kumar. Integrating scientific knowledge with machine learning for engineering and environmental systems, 2022.
- [26] Charlie Wood. Powerful ‘ machine scientists ’ distill the laws of physics from raw data. *Quanta Magazine*, 2022.
- [27] Tony Worm and Kenneth Chiu. Scaling up prioritized grammar enumeration for scientific discovery in the cloud. In *Proceedings - 2014 IEEE International Conference on Big Data, IEEE Big Data 2014*, 2014.

## 7 Project Management:

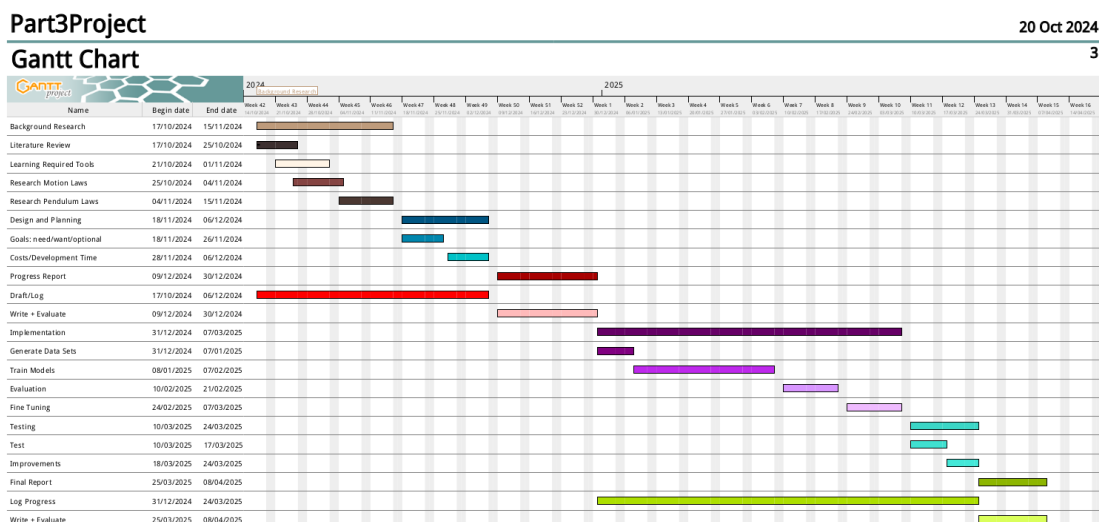
### 7.1 Risk Assessment:

Prior to the commencement of this project, a risk assessment was conducted to anticipate potential challenges and plan mitigation strategies. This proactive approach aimed to ensure a smoother development process. Several risks did materialize during the course of the work; however, thanks to the preparations made, they were effectively managed and did not significantly hinder progress.

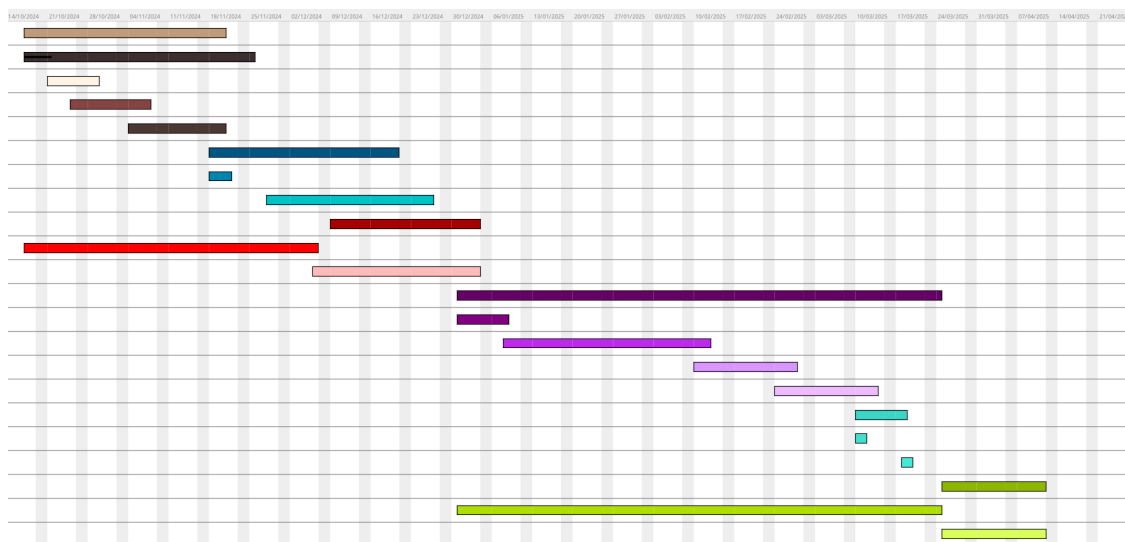
Issue	Impact	Prob	Risk	Mitigation
Unexpected delays and accidents	3	3	7	Include contingency plans and a 3 week break between major stages of the project, to allow for unexpected incidents.of the project, to allow for unexpected incidents.
Unable to generate enough experimental data due to lack of computational power.	4	1	14	Explore alternate more efficient ways of simulating data, consider using cloud infrastructure or potentially the Universities HPC facilities.
Challenges learning the pendulum laws and the derivation	2	4	5	Seek other resources from the Physics Department to learn the Physics required. Look up explanations online to learn.
Interpretability Challenges	3	2	10	Challenges in interpreting how the model works, can be mitigated through visualising the data, plotting results and through seeking ways to explain the model.
Symbolic Regression Failure	3	4	9	Research and select appropriate SR algorithms/libraries; Start with a simple operator set; Test on synthetic data with known solutions.
Scope Creep	2	2	14	Define clear project scope and deliverables at the start; Use a detailed timeline with milestones; Regularly review progress against the plan.

## 7.2 Gantt Charts:

A Gantt chart along with a rough outline of the relevant dates for various submission was made towards the beginning of this project, this also included contingency planning and short yet frequent breaks every couple weeks. Furthermore also made a gantt chart to compare the actual timeline vs the prediction made.



What actually ended up happening:



## 7.3 Project Reflection:

In reflection, the symbolic regressor was successful and has demonstrated all of the goals initially outlined, successfully deriving the expected formulas. However, one extension that was intended—deriving a chaotic system like a double pendulum—proved to be extremely difficult and ultimately impossible with the current architecture. Even after removing the dimensions from the data and eliminating constants, the sheer number of variables involved in predicting angular momentum, along with the fact that it is a coupled system, made the regressor's design inadequate. Even when the dataset was split and the data fed in individually, the search

space remained large, and valid expressions only emerged at a recursive depth of 5. This depth led to a total of 254,803,968 possible searches—an amount which, by calculation, would take longer than the age of the universe to solve using the current algorithm.




Figure 9: (a) Static double pendulum




Figure 10: (b) Double pendulum path

Figure 11: Side-by-side visualisation of a double pendulum's static setup and path.

Having never previously worked on a project of this academic scale, this was a novel experience. Every decision and design choice had to be justified with prior research and academic references, making the process significantly more demanding than expected. Justifying the architectural choices and constantly being under pressure to produce valid expressions drove the development of the regressor in the most effective way possible. A Golang version of the regressor was briefly explored but wisely abandoned due to time constraints, allowing focus to return to writing the report. Some aspects of the project took far longer than anticipated, and attempts to collaborate with authors of referenced papers were unsuccessful, as no responses were received.

An agile approach would not have worked well in this context. It proved more valuable to focus on the overall design, plan for the long term, and rigorously test components throughout, rather than developing a half-baked and ineffective core. This rationale guided the early stages, which involved careful library selection and thorough research into various strategies before committing to a design. Each function was rigorously tested upon writing to ensure speed, effectiveness, and correctness. Modules were developed in such a way that, when combined, the main regressor worked with minimal or no significant bugs. Furthermore, planning contingency time in the form of short but frequent breaks proved highly beneficial and was far more effective than reserving all available time until the end.

While implementing, designing, testing, and writing the project and report was challenging, the process led to the development of several key research skills—such as referencing properly, reading academic papers more systematically, and learning that seemingly minor details can cause the longest delays. It became clear that it is ambitious and often unwise to expect to achieve absolutely everything envisioned at the start of a project. It is always better to err on the side of caution—promise little and overdeliver, rather than the opposite. Working at the intersection of machine learning and its applications in other domains such as biology and physics has highlighted the immense potential of machine learning. It has shown how tools based on algorithms and computation can solve problems that once required exceptional human insight—now solved by silicon and code.

## 8 Conclusion:

To summarise the goal of this project is to explore deriving physical laws through AI, and helping automate the discovery of empirical laws. Hopefully, this would help speed up the process of discovering new laws, in data, as well as remove the inherent human bias towards creating expressions. A significant amount of attention was placed on reimplementing the regressor from scratch in a more modern and maintainable language, in order to make it easier for future improvements to be made. Also I've used the least amount of libraries possible (numpy and sympy), in order to build this, allowing it to be really easy to maintain the codebase. It has successfully derived the resulted expected and has generated a series of complex expressions for describing data.

The Symbolic Regressor can also utilise multiple different methods, such as dimensional analysis, neural networks to help visual the problem, pareto frontiers to optimise between accuracy and complexity, in order to approach the same dataset in multiple ways, not only helping increase accuracy, however there seems to be some edge cases that causes errors due to mismatched types, in some more complex non-linear datasets. Python's dynamic type system makes it difficult to debug this error.

During the chaining of the variety of methods, I had realised that the gradient and prediction information from the neural network wasn't used to its maximum potential, and this aspect could be further improved to then implement approximations or approach using neural networks in a slightly different manner, using it in small batch learning instead. This strikes a balance between accuracy and interpretability.

Recommendation for future works, is to extend the use of neural networks and utilise the produced data in better visualisation methods besides just a simple plot. The polynomial fit module can be further extended to give an approximation when a brute force or dimensional analysis approach doesn't prove to be fruitful. In practise this architecture is far faster in deducing the expressions from a much smaller set of data than the original aifeynman model, as it uses minimal libraries (mostly only numpy and sympy), and is therefore far faster, hours vs minutes sometimes seconds to solve the same expression. But in terms of accuracy my model suffers due to the affront mentioned type mismatch errors. Furthermore I have started working on a golang version, which has potential to be far faster than PySR, as goroutines inherently offer better optimisations and speed, now if we use the hybrid approach found in [4], modelling the sequence as go structs offers this advantage, potentially increasing the speed hundreds of times. I tried to brute force for a double pendulum equation, with some minimal filtering methods, however the search space was too vast due to the chaotic nature of the problem as well as the large number of variables involved even when simplifying the expression.

There is also potential to explore this beyond physics and biological data, chemistry for example is another natural step that can be taken, but this can be applied in a broader sense to any data, as long as it is coded correctly, in a way the model can understand, giving predictions on events based on data, ie potential economic data, really any "variables" that have causation and correlation with each other.

future works

genetic programming can help prune even more, and is useful for approximations, trainable meijer g functions , (Crabbe et al. 2020), (Alaa van der Schaar 2019 Genetic programming (GP) has traditionally been a core technique for symbolic regression [references...], although challenges related to scaling and overfitting have been noted (see discussion in Abdellaoui Mehrkanon, 2021, citing [24]), leading to the exploration of alternative methods.